

Module Guide for Software Engineering

Team 21, Visionaries
Angela Zeng
Ann Shi
Ibrahim Sahi
Manan Sharma
Stanley Chen

November 13, 2025

1 Revision History

Date	Version	Notes
Date 1	1.0	Notes
Date 2	1.1	Notes

2 Reference Material

This section records information for easy reference.

2.1 Abbreviations and Acronyms

symbol	description
AC	Anticipated Change
DAG	Directed Acyclic Graph
M	Module
MG	Module Guide
OS	Operating System
R	Requirement
SC	Scientific Computing
SRS	Software Requirements Specification
Software Engineering	Explanation of program name
UC	Unlikely Change
[etc. —SS]	[... —SS]

Contents

1 Revision History	i
2 Reference Material	ii
2.1 Abbreviations and Acronyms	ii
3 Introduction	1
4 Anticipated and Unlikely Changes	2
4.1 Anticipated Changes	2
4.2 Unlikely Changes	2
5 Module Hierarchy	2
6 Connection Between Requirements and Design	3
7 Module Decomposition	5
7.1 Hardware Hiding Module (M1)	5
7.2 Behaviour-Hiding Modules	6
7.2.1 Data Ingestion Module (M2)	6
7.2.2 Real-Time Streaming Module (M3)	6
7.2.3 Dashboard Visualization Module (M4)	7
7.2.4 Reporting Module (M5)	7
7.3 Software Decision Modules	8
7.3.1 Data Preprocessing Module (M6)	8
7.3.2 Privacy & Infrastructure Module (M7)	8
7.3.3 Engagement Analytics Module (M8)	9
7.3.4 Correlation & Visual Analysis Module (M9)	9
8 Traceability Matrix	9
9 Use Hierarchy Between Modules	10
10 User Interfaces	11
11 Design of Communication Protocols	11
12 Timeline	11

List of Tables

1 Module Hierarchy	3
2 Trace Between Requirements and Modules	10

3	Trace Between Anticipated Changes and Modules	10
---	---	----

List of Figures

1	Use hierarchy among modules	11
---	---------------------------------------	----

3 Introduction

Decomposing a system into modules is a commonly accepted approach to developing software. A module is a work assignment for a programmer or programming team (Parnas et al., 1984). We advocate a decomposition based on the principle of information hiding (Parnas, 1972). This principle supports design for change, because the “secrets” that each module hides represent likely future changes. Design for change is valuable in SC, where modifications are frequent, especially during initial development as the solution space is explored.

Our design follows the rules laid out by Parnas et al. (1984), as follows:

- System details that are likely to change independently should be the secrets of separate modules.
- Each data structure is implemented in only one module.
- Any other program that requires information stored in a module’s data structures must obtain it by calling access programs belonging to that module.

After completing the first stage of the design, the Software Requirements Specification (SRS), the Module Guide (MG) is developed (Parnas et al., 1984). The MG specifies the modular structure of the system and is intended to allow both designers and maintainers to easily identify the parts of the software. The potential readers of this document are as follows:

- New project members: This document can be a guide for a new project member to easily understand the overall structure and quickly find the relevant modules they are searching for.
- Maintainers: The hierarchical structure of the module guide improves the maintainers’ understanding when they need to make changes to the system. It is important for a maintainer to update the relevant sections of the document after changes have been made.
- Designers: Once the module guide has been written, it can be used to check for consistency, feasibility, and flexibility. Designers can verify the system in various ways, such as consistency among modules, feasibility of the decomposition, and flexibility of the design.

The rest of the document is organized as follows. Section 4 lists the anticipated and unlikely changes of the software requirements. Section 5 summarizes the module decomposition that was constructed according to the likely changes. Section 6 specifies the connections between the software requirements and the modules. Section 7 gives a detailed description of the modules. Section 8 includes two traceability matrices. One checks the completeness of the design against the requirements provided in the SRS. The other shows the relation between anticipated changes and the modules. Section 9 describes the use relation between modules.

4 Anticipated and Unlikely Changes

This section lists possible changes to the system. According to the likeliness of the change, the possible changes are classified into two categories. Anticipated changes are listed in Section 4.1, and unlikely changes are listed in Section 4.2.

4.1 Anticipated Changes

Anticipated changes are the source of the information that is to be hidden inside the modules. Ideally, changing one of the anticipated changes will only require changing the one module that hides the associated decision. The approach adapted here is called design for change.

AC1: The specific hardware on which the software is running.

AC2: The format of the initial input data.

...

[Anticipated changes relate to changes that would be made in requirements, design or implementation choices. They are not related to changes that are made at run-time, like the values of parameters. —SS]

4.2 Unlikely Changes

The module design should be as general as possible. However, a general system is more complex. Sometimes this complexity is not necessary. Fixing some design decisions at the system architecture stage can simplify the software design. If these decision should later need to be changed, then many parts of the design will potentially need to be modified. Hence, it is not intended that these decisions will be changed.

UC1: Input/Output devices (Input: File and/or Keyboard, Output: File, Memory, and/or Screen).

...

5 Module Hierarchy

This section provides an overview of the module design. Modules are summarized in a hierarchy decomposed by secrets in Table 1. The modules listed below, which are leaves in the hierarchy tree, are the modules that will actually be implemented.

M1: Hardware-Hiding Module

M2: Data Ingestion Module

- M3:** Real-Time Streaming Module
- M4:** Dashboard Visualization Module
- M5:** Reporting Module
- M6:** Data Preprocessing Module
- M7:** Privacy & Infrastructure Module
- M8:** Engagement Analytics Module
- M9:** Correlation & Visual Analysis Module

Level 1	Level 2
Hardware-Hiding Module	Hardware-Hiding Module (M1)
Behaviour-Hiding Module	Data Ingestion Module (M2) Real-Time Streaming Module (M3) Dashboard Visualization Module (M4) Reporting Module (M5)
Software Decision Module	Data Preprocessing Module (M6) Privacy & Infrastructure Module (M7) Engagement Analytics Module (M8) Correlation & Visual Analysis Module (M9)

Table 1: Module Hierarchy

6 Connection Between Requirements and Design

The design of the system is intended to satisfy the requirements developed in the SRS. In this stage, the system is decomposed into modules. The connection between requirements and modules is listed in Table 2.

[The intention of this section is to document decisions that are made “between” the requirements and the design. To satisfy some requirements, design decisions need to be made. Rather than make these decisions implicit, they are explicitly recorded here. For instance, if a program has security requirements, a specific design decision may be made to satisfy those requirements with a password. —SS]

- **Separation of data ingestion from preprocessing:** Requirements relating to multi-device input, raw data capture, timing synchronization, and recovery from connection loss motivated the creation of a dedicated *Data Ingestion Module*. Requirements concerning noise reduction, calibration correction, homography projection, and coordinate normalization are distinct and subject to change independently, motivating the *Data Preprocessing Module*. This reflects the SRS separation between raw data capture and analytic data preparation.
- **Explicit handling of real-time constraints:** The need for live metrics, dashboard updating, latency guarantees, and real-time streaming (as described in the SRS operational modes) led to the creation of a separate *Real-Time Streaming Module*. This encapsulates decisions about streaming frameworks, buffering policy, and delivery latency, which are likely to evolve as performance constraints change.
- **Distinguishing analytics computation from visualization:** Engagement metrics, instructor–student gaze alignment, heatmap generation, and correlation analyses are internal computations not directly described to the end user. These behaviours motivated two separate software-decision modules: the *Engagement Analytics Module* and the *Correlation & Visual Analysis Module*. Meanwhile, all visible UI behaviours—including real-time views, playback, filtering, and seat-map interactions—are grouped in the *Dashboard Visualization Module*. This follows the SRS distinction between analytic computation and user-facing presentation.
- **Privacy and anonymization as a first-class design concern:** The SRS explicitly lists privacy, anonymization, secure handling of raw video, and compliance with data retention constraints. These requirements cannot be embedded inside ingestion, analytics, or dashboard logic without creating strong coupling. They therefore motivate the independent *Privacy & Infrastructure Module*, which encapsulates the design decisions related to face blurring, spatial masking, secure storage, authorization, and logging.
- **Support for offline summaries and instructor reporting:** Requirements calling for post-session summaries, exportable metrics, and instructor-friendly performance reports lead to the *Reporting Module*. This module encapsulates decisions about report structure, export format, and summarization logic—decisions that are likely to change independently of both preprocessing and dashboard design.
- **Abstraction of physical devices and OS services:** The SRS assumes multiple hardware sources (eye-tracking glasses, central camera, audio, file storage, and network interfaces). To prevent downstream modules from depending on specific device drivers or OS-level behaviour, a general *Hardware-Hiding Module* is included. This encapsulates the design decision to virtualize camera and sensor access, a key anticipated change identified in the hazards and SRS (e.g., replacing Neon with another device).

- **Backend service and session-lifecycle requirements:** The SRS specifies requirements for creating, managing, and retrieving recorded sessions, as well as backend API access for live or historical data (SRS: Service & API requirements; Session Lifecycle Requirements). These behaviours are not purely real-time streaming nor visualization concerns. This motivated folding backend session-control logic into the *Real-Time Streaming Module*, which now encapsulates API endpoints for session creation, playback control, and client subscriptions.
- **Secure storage, retention, and access-control requirements:** The SRS defines requirements for data retention policies, authenticated access, role-based controls (e.g., instructor vs. researcher), and logging of data access. Embedding these into ingestion or analytics would create tight coupling, so these decisions are isolated in the *Privacy & Infrastructure Module*. This module therefore owns secure storage, retention strategy, authorization, and audit logging as design decisions hidden from the rest of the system.

Overall, the module structure shows the boundaries between behaviours visible to the user, performance-driven internal computations, privacy-critical infrastructure, and hardware-specific interactions. This makes sure that changes in requirements - such as improved analytics, new privacy constraints, or new hardware - affect only the module that hides the corresponding decision.

7 Module Decomposition

Modules are decomposed according to the principle of “information hiding” proposed by [Parnas et al. \(1984\)](#). The *Secrets* field in a module decomposition is a brief statement of the design decision hidden by the module. The *Services* field specifies *what* the module will do without documenting *how* to do it. For each module, a suggestion for the implementing software is given under the *Implemented By* title. If the entry is *OS*, this means that the module is provided by the operating system or by standard programming language libraries. *Software Engineering* means the module will be implemented by the Software Engineering software.

Only the leaf modules in the hierarchy have to be implemented. If a dash (-) is shown, this means that the module is not a leaf and will not have to be implemented.

7.1 Hardware Hiding Module (M1)

Secrets: The device-specific interfaces, drivers, and low-level OS mechanisms used to access eye-tracking glasses (Neon), world camera input, microphones, GPU acceleration, file system APIs, local and cloud storage backends, network sockets, and system time synchronization (e.g., NTP services).

Services: Provides a virtualized hardware interface for capturing gaze streams, world camera frames, audio signals, timestamps, and device metadata, and for reading/writing

data to persistent storage. Exposes uniform APIs so downstream modules do not depend on specific device models, storage engines, or operating systems.

Implemented By: OS, vendor SDKs, device drivers

Type of Module: Library

7.2 Behaviour-Hiding Modules

Secrets: The externally visible behaviours of the system, including session playback, dashboard display logic, ingestion behaviour, and reporting rules as defined in the SRS.

Services: Provides all user-facing behaviours, connecting virtualized hardware to internal algorithms. These modules must change if the SRS behaviour changes.

Implemented By: –

7.2.1 Data Ingestion Module (M2)

Secrets: The method of connecting to and receiving data from multiple hardware sources (glasses, cameras, microphones), including device protocols, buffering policy, timestamp alignment strategy (including use of system/NTP time), and fault-recovery behaviour for dropped connections or device failures.

Services:

- Accepts raw gaze, video, audio, and metadata streams from all devices.
- Performs initial timestamping and session indexing using a common time base.
- Manages session start/stop hooks that allow backend services to create sessions, register participants, and mark completed recordings.
- Outputs unified raw data packets for preprocessing and optional storage.

Implemented By: Software Engineering

Type of Module: Abstract Object

7.2.2 Real-Time Streaming Module (M3)

Secrets: The design decisions for real-time transport (Kafka, WebSockets), latency buffering, stream synchronization, retry/timeout strategies, throttling, and the external API endpoints through which clients subscribe to streams and control session state (e.g., REST/GraphQL/WebSocket interfaces).

Services:

- Exposes backend APIs for creating, listing, and controlling live or replay sessions (start, pause, stop) used by the dashboard and other clients.
- Delivers live preprocessed data and analytics to the dashboard and other authorized consumers with bounded latency.

- Maintains real-time session state (live/offline toggle, paused state, health status) and reports it to monitoring/logging infrastructure.
- Ensures smooth playback for recorded sessions by streaming stored data through the same interfaces used for live sessions.

Implemented By: Software Engineering

Type of Module: Abstract Object

7.2.3 Dashboard Visualization Module (M4)

Secrets: The visualization grammar, UI layout, interaction rules, playback controls, and display logic for heatmaps, timelines, seat maps, and instructor–student alignment graphs.

Services:

- Presents real-time and historical analytics.
- Renders heatmaps, fixation charts, engagement over time, and gaze-alignment plots.
- Supports filters (student, seat zone, activity type, time window).
- Provides playback and zoom-in on individual student streams.

Implemented By: Software Engineering

Type of Module: Abstract Object

7.2.4 Reporting Module (M5)

Secrets: The structure and formatting of summaries, metrics, annotations, and export formats (PDF/HTML) and the rules for instructor reflection.

Services:

- Generates session summaries, attention-duration metrics, and engagement timelines.
- Produces exportable reports for instructors and teaching portfolios.
- Supports annotation, note-taking, screengrabs, and timeline cropping.

Implemented By: Software Engineering

Type of Module: Library

7.3 Software Decision Modules

Secrets: Internal algorithms, mathematical models, data structures, and analytic methods that are not externally visible. These change primarily due to performance or research-driven improvements, not due to changes in behaviour.

Services: Provides internal data processing, analytics, correlations, and privacy enforcement services for behaviour-hiding modules.

Implemented By: –

7.3.1 Data Preprocessing Module (M6)

Secrets: Filtering algorithms (smoothing, denoising), calibration transformations, gaze-to-screen homography methods, fixation detection rules, and temporal alignment techniques.

Services:

- Converts raw gaze vectors into screen/world-space coordinates.
- Denoises and filters gaze data.
- Identifies fixations, saccades, attention windows.
- Standardizes multimodal streams for analytics.

Implemented By: Software Engineering

Type of Module: Abstract Data Type

7.3.2 Privacy & Infrastructure Module (M7)

Secrets: The anonymization pipeline, face-blurring method, device obfuscation, consent rules, access-control and role-based authorization policies, logging and monitoring configuration, database schema and storage layout for raw and processed data, and data retention and deletion strategy.

Services:

- Performs face blurring, device masking, and area-of-interest redaction on video and gaze data before it is exposed to user-facing modules.
- Ensures audio and video privacy filtering before data reaches the dashboard.
- Manages authentication and authorization for backend APIs and the dashboard (e.g., instructor vs. researcher roles).
- Provides secure storage and retrieval of raw and processed data in the underlying database or file system according to defined retention rules.
- Provides logging and monitoring hooks (e.g., audit logs, API access logs, health checks) used by other modules.

Implemented By: Software Engineering

Type of Module: Library

7.3.3 Engagement Analytics Module (M8)

Secrets: The mathematical models used to compute engagement metrics, moving averages, attention peaks/troughs, activity-phase detection, and time-on-task calculations.

Services:

- Computes engagement over time for individuals and aggregates.
- Detects disengagement periods and high-attention intervals.
- Classifies lecture activities (discussion, slides, group work).
- Generates metrics for reporting and dashboard visualization.

Implemented By: Software Engineering

Type of Module: Abstract Data Type

7.3.4 Correlation & Visual Analysis Module (M9)

Secrets: The gaze-correlation model, instructor–student alignment formulas, visual material detection (text block, image region), and methods for transforming heatmaps into interpretable coordinate-space analytics.

Services:

- Computes correlations between instructor gaze and class gaze.
- Maps class attention onto slide regions or world-camera space.
- Identifies which visual elements students focus on, and how this changes over time.
- Supports seat-map and spatial-zone based metrics and filters.

Implemented By: Software Engineering

Type of Module: Abstract Data Type

8 Traceability Matrix

This section shows two traceability matrices: between the modules and the requirements and between the modules and the anticipated changes.

Req.	Modules
R1	M1, M??, M??, M??
R2	M??, M??
R3	M??
R4	M??, M??
R5	M??, M??, M??, M??, M??, M??
R6	M??, M??, M??, M??, M??, M??
R7	M??, M??, M??, M??, M??
R8	M??, M??, M??, M??, M??
R9	M??
R10	M??, M??, M??
R11	M??, M??, M??, M??

Table 2: Trace Between Requirements and Modules

AC	Modules
AC1	M1
AC2	M??
AC??	M??

Table 3: Trace Between Anticipated Changes and Modules

9 Use Hierarchy Between Modules

In this section, the uses hierarchy between modules is provided. Parnas (1978) said of two programs A and B that A *uses* B if correct execution of B may be necessary for A to complete

the task described in its specification. That is, A *uses* B if there exist situations in which the correct functioning of A depends upon the availability of a correct implementation of B. Figure 1 illustrates the use relation between the modules. It can be seen that the graph is a directed acyclic graph (DAG). Each level of the hierarchy offers a testable and usable subset of the system, and modules in the higher level of the hierarchy are essentially simpler because they use modules from the lower levels.

[The uses relation is not a data flow diagram. In the code there will often be an import statement in module A when it directly uses module B. Module B provides the services that module A needs. The code for module A needs to be able to see these services (hence the import statement). Since the uses relation is transitive, there is a use relation without an import, but the arrows in the diagram typically correspond to the presence of import statement. —SS]

[If module A uses module B, the arrow is directed from A to B. —SS]

Figure 1: Use hierarchy among modules

10 User Interfaces

[Design of user interface for software and hardware. Attach an appendix if needed. Drawings, Sketches, Figma —SS]

11 Design of Communication Protocols

[If appropriate —SS]

12 Timeline

[Schedule of tasks and who is responsible —SS]

[You can point to GitHub if this information is included there —SS]

References

David L. Parnas. On the criteria to be used in decomposing systems into modules. *Comm. ACM*, 15(2):1053–1058, December 1972.

David L. Parnas. Designing software for ease of extension and contraction. In *ICSE '78: Proceedings of the 3rd international conference on Software engineering*, pages 264–277, Piscataway, NJ, USA, 1978. IEEE Press. ISBN none.

D.L. Parnas, P.C. Clement, and D. M. Weiss. The modular structure of complex systems. In *International Conference on Software Engineering*, pages 408–419, 1984.